





CartPole learning and retention

Partial GA/ES reporting; PPO deferred · 8/20 trials

GA · $n = 4$

ES · $n = 4$



— Own task - - - Previous task

Thin: individual trials; bold: arithmetic mean. Native update clocks; no smoothing.

Checkpoint agents use ten fresh episodes per target; previous-task return starts at phase 2.